
PERFORMANCE TUNING

- We add indexes , primary keys.
- We create table partitions.
- We analyze query execution plan.
- Remove unnecessary large-table full-table scans
- Cache small-table full-table scans
- Verify optimal index usage
- Using hints to tune Oracle SQL
- Self-order the table joins

PERFORMANCE TUNING

NOTHING !!!!!

PERFORMANCE TUNING

- How sensibly you use this tool.
- Which helps to reduce processing cost.
- Which helps to reduce storage cost.

PERFORMANCE TUNING

- `SELECT * FROM EMP;`
- `SELECT EMP_NAME, EMP_ADDR FROM EMP;`
- Sharing virtual warehouse while dev activities.
- Order your data by filter columns during data loading process.
- Use multi cluster warehouse instead of spinning up existing cluster to bigger size.

PERFORMANCE TUNING

- There is no concept of index.
- No concept of primary key , foreign key constraints.
- No need of transaction management.
- There is no buffer pool.
- You will never encounter out of memory exception.

PERFORMANCE TUNING

- But How ACID transactions are possible !!!!

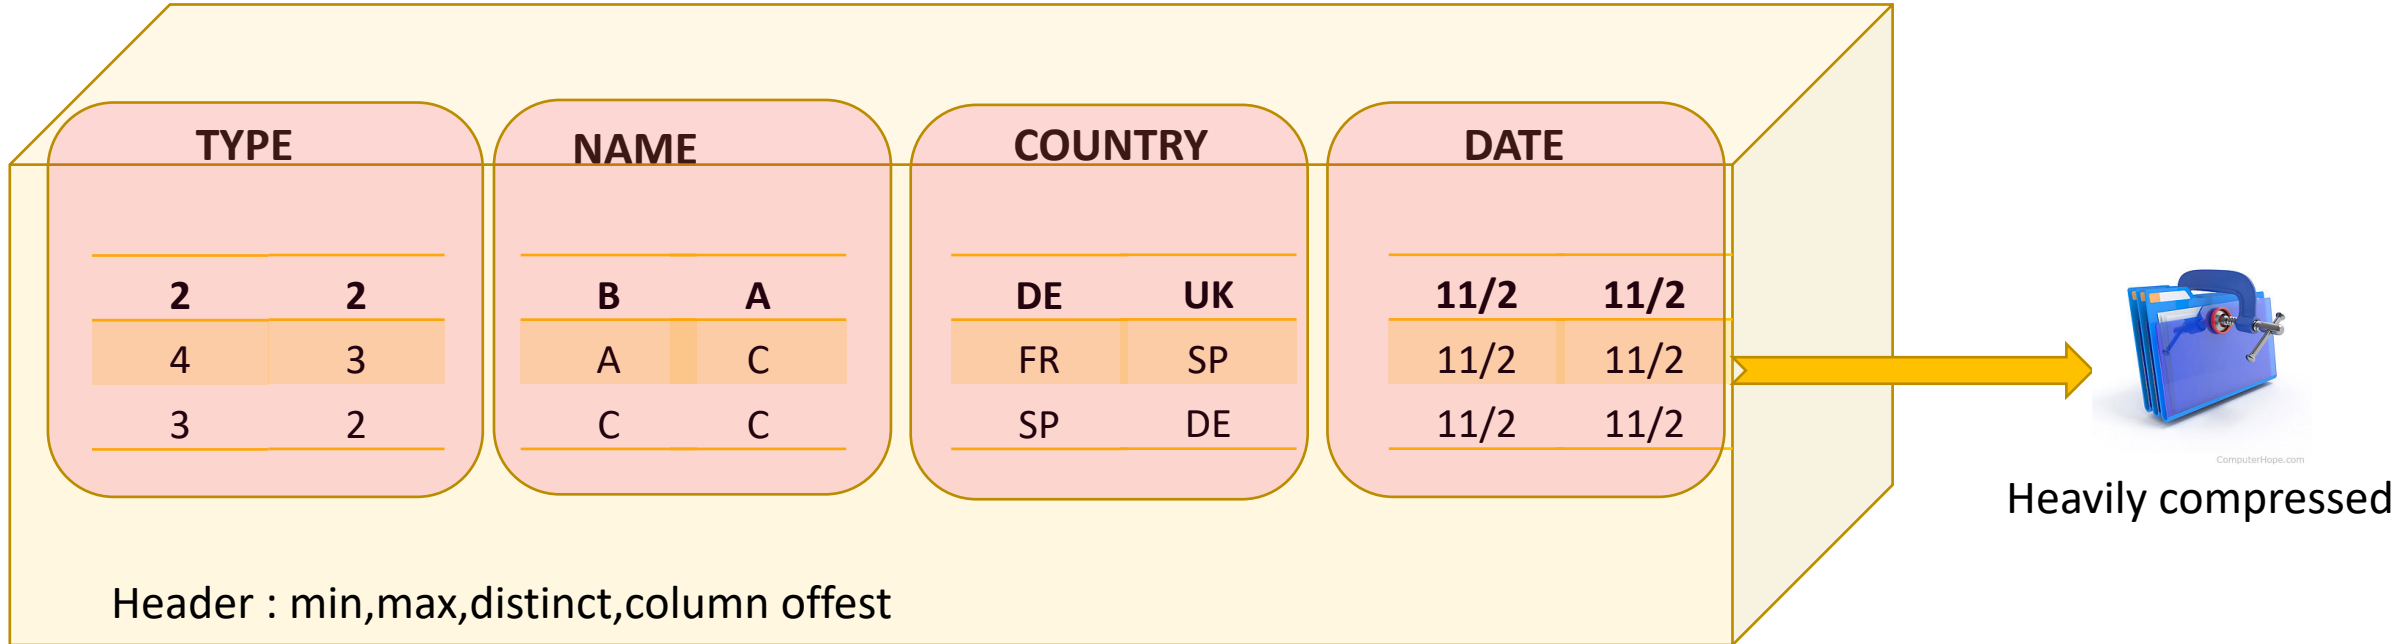
Query management and optimization

- Parsing.
- Object resolution.
- Access control.
- Plan optimization.
- Snowflake will not use indexes.

Query management and optimization

- Storage medium in snowflake is s3 and data format is compressed files.
- Maintaining indexes significantly increases volume of data and data loading time.
- User need to explicitly create indices. Which goes against snowflake philosophy of SAS.
- Maintaining indices can be complex , expensive and risky process.

Table File



Micro-partitioning is automatically **performed** on all Snowflake tables. Tables are transparently partitioned using the **ordering of the data as it is inserted/loaded**.

Query management and optimization

- Pruning.
- Zone maps.
- Data skipping.


Concurrency control

KEY	VALUE
TABLE1	S3:\\<some_location>\
TABLE1	S3:\\<some_location>\
TABLE1	S3:\\<some_location>\
TABLE1	S3:\\<some_location>\

Concurrency control



Concurrency control



SNAPSHOT ISOLATION

HOW UPDATES WORKS

- DELETE
- INSERT

HOW UPDATES WORKS

Key	Value	Active
EMP	1	Y
EMP	2	N
EMP	3	Y
EMP	4	Y



Marked as deleted



Inserted

HOW UPDATES WORKS

Partitioned table and DML

- No update-in-places (AWS S3, Azure Blob are immutable file systems)
- DML **adds** or **deletes** partition files -> partition files are written in full

DML 1

Copy into t from @stage1

ID	Name	
1	John	P 1
2	Scott	
3	Mary	
4	Jane	P 2
5	Jack	
6	Claire	

DML 1: + P 1
+ P 2

DML 2

Insert into t values (7, 'Pierre')

ID	Name	
1	John	P 1
2	Scott	
3	Mary	
4	Jane	P 2
5	Jack	
6	Claire	
7	Pierre	P 3

DML 2: + P 3

DML 3

Delete from t where name = 'Jane'

ID	Name	
1	John	P 1
2	Scott	
3	Mary	
5	Jack	P 4
6	Claire	
7	Pierre	P 3

DML 3: - P 2
+ P 4

Why I have to remove files

- Because we are using blob storage area like s3. Here you can only overwrite full files.

LEASON **LEARNED**

- Before executing update, statement check how many records will get impacted. If more than 80% of records are getting impacted, then instead of executing update you can consider recreating whole table.
- You can execute delete and insert statement separately.
- When you are trying to update or delete try considering numeric columns as its easy to scan micro- partitions.